

# CS7643: Deep Learning

## Spring 2020

### Problem Set 1

Instructor: Zsolt Kira

TAs: Yihao Chen, Sameer Dharur, Rahul Duggal, Patrick Grady, Harish Kamath  
Yinquan Lu, Anishi Mehta, Manas Sahni, Jiachen Yang, Zhuoran Yu

Discussions: <https://piazza.com/gatech/spring2020/cs4803dl7643a/home>

Due: Tuesday, February 11, 11:55pm

#### Instructions

1. We will be using Gradescope to collect your assignments. Please read the following instructions for submitting to Gradescope carefully!
  - Each subproblem must be submitted on a separate page. When submitting to Gradescope, make sure to mark which page(s) corresponds to each problem/sub-problem. For instance, Q5 has 5 subproblems, and the solution to each must start on a new page. Similarly, Q8 has 8 subproblems, and the writeup for each should start on a new page.
  - For the coding problems (Q8), please use the provided `collect_submission.sh` script and upload `hw1.zip` to the HW1 Code assignment on Gradescope. While we will not be explicitly grading your code, you are still required to submit it. Please make sure you have saved the most recent version of your jupyter notebook before running this script. Further, append the writeup for each Q8 subproblem to your PS1 solution PDF.
  - Note: This is a large class and Gradescope's assignment segmentation features are essential. Failure to follow these instructions may result in parts of your assignment not being graded. We will not entertain regrading requests for failure to follow instructions.
2. L<sup>A</sup>T<sub>E</sub>X'd solutions are strongly encouraged (solution template available at [cc.gatech.edu/classes/AY2020/cs7643\\_fall/assets/sol1.tex](http://cc.gatech.edu/classes/AY2020/cs7643_fall/assets/sol1.tex)), but scanned handwritten copies are acceptable. Hard copies are **not** accepted.
3. We generally encourage you to collaborate with other students.

You may talk to a friend, discuss the questions and potential directions for solving them. However, you need to write your own solutions and code separately, and *not* as a group activity. Please list the students you collaborated with.

## 1 Gradient Descent

1. (3 points) We often use iterative optimization algorithms such as Gradient Descent to find  $\mathbf{w}$  that minimizes a loss function  $f(\mathbf{w})$ . Recall that in gradient descent, we start with an initial

value of  $\mathbf{w}$  (say  $\mathbf{w}^{(1)}$ ) and iteratively take a step in the direction of the negative of the gradient of the objective function *i.e.*

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \nabla f(\mathbf{w}^{(t)}) \quad (1)$$

for learning rate  $\eta > 0$ .

In this question, we will develop a slightly deeper understanding of this update rule, in particular for minimizing a convex function  $f(\mathbf{w})$ . Note: this analysis will not directly carry over to training neural networks since loss functions for training neural networks are typically not convex, but this will (a) develop intuition and (b) provide a starting point for research in non-convex optimization (which is beyond the scope of this class).

Recall the first-order Taylor approximation of  $f$  at  $\mathbf{w}^{(t)}$ :

$$f(\mathbf{w}) \approx f(\mathbf{w}^{(t)}) + \langle \mathbf{w} - \mathbf{w}^{(t)}, \nabla f(\mathbf{w}^{(t)}) \rangle \quad (2)$$

When  $f$  is convex, this approximation forms a lower bound of  $f$ , *i.e.*

$$f(\mathbf{w}) \geq \underbrace{f(\mathbf{w}^{(t)}) + \langle \mathbf{w} - \mathbf{w}^{(t)}, \nabla f(\mathbf{w}^{(t)}) \rangle}_{\text{affine lower bound to } f(\cdot)} \quad \forall \mathbf{w} \quad (3)$$

Since this approximation is a ‘simpler’ function than  $f(\cdot)$ , we could consider minimizing the approximation instead of  $f(\cdot)$ . Two immediate problems: (1) the approximation is affine (thus unbounded from below) and (2) the approximation is faithful for  $\mathbf{w}$  close to  $\mathbf{w}^{(t)}$ . To solve both problems, we add a squared  $\ell_2$  *proximity term* to the approximation minimization:

$$\underset{\mathbf{w}}{\operatorname{argmin}} \underbrace{f(\mathbf{w}^{(t)}) + \langle \mathbf{w} - \mathbf{w}^{(t)}, \nabla f(\mathbf{w}^{(t)}) \rangle}_{\text{affine lower bound to } f(\cdot)} + \underbrace{\frac{\lambda}{2} \left\| \mathbf{w} - \mathbf{w}^{(t)} \right\|^2}_{\text{trade-off proximity term}} \quad (4)$$

Notice that the optimization problem above is an unconstrained quadratic programming problem, meaning that it can be solved in closed form (hint: gradients).

What is the solution  $\mathbf{w}^*$  of the above optimization? What does that tell you about the gradient descent update rule? What is the relationship between  $\lambda$  and  $\eta$ ?

Answer: Take the derivative of  $w$ :

$$f'(w) = \nabla f(\mathbf{w}^{(t)}) + \lambda(\mathbf{w} - \mathbf{w}^{(t)}) = 0 \quad (5)$$

$$\mathbf{w}^* = \mathbf{w}^{(t)} - \frac{1}{\lambda} \nabla f(\mathbf{w}^{(t)}) \quad (6)$$

$$\eta = \frac{1}{\lambda} \quad (7)$$

Gradient decent can minimize the approximation function at each step. The value of  $\lambda$  means the penalty of proximity term. When  $\lambda$  increases, the learning rate will decrease, which means the step along GD can be small, and vice versa.

2. (3 points) Let’s prove a lemma that will initially seem devoid of the rest of the analysis but will come in handy in the next sub-question when we start combining things. Specifically, the analysis in this sub-question holds for any  $\mathbf{w}^*$ , but in the next sub-question we will use it for  $\mathbf{w}^*$  that minimizes  $f(\mathbf{w})$ .

Consider a sequence of vectors  $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_T$ , and an update equation of the form  $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \mathbf{v}_t$  with  $\mathbf{w}^{(1)} = \mathbf{0}$ . Show that:

$$\sum_{t=1}^T \langle \mathbf{w}^{(t)} - \mathbf{w}^*, \mathbf{v}_t \rangle \leq \frac{\|\mathbf{w}^*\|^2}{2\eta} + \frac{\eta}{2} \sum_{t=1}^T \|\mathbf{v}_t\|^2 \quad (8)$$

Answer:

$$\langle \mathbf{w}^{(t)} - \mathbf{w}^*, \mathbf{v}_t \rangle = \frac{1}{\eta} \langle \mathbf{w}^{(t)} - \mathbf{w}^*, \eta \mathbf{v}_t \rangle \quad (9)$$

$$= \left\| \mathbf{w}^{(t)} - \mathbf{w}^* \right\|^2 + \|\eta \mathbf{v}_t\|^2 + \left\| \mathbf{w}^{(t)} - \mathbf{w}^* - \eta \mathbf{v}_t \right\|^2 \quad (10)$$

$$= \frac{1}{4\eta} \left( \left\| \mathbf{w}^{(t)} - \mathbf{w}^* + \eta \mathbf{v}_t \right\|^2 - \left\| \mathbf{w}^{(t)} - \mathbf{w}^* - \eta \mathbf{v}_t \right\|^2 \right) \quad (11)$$

$$= \frac{1}{2\eta} \left( \left\| \mathbf{w}^{(t)} - \mathbf{w}^* \right\|^2 - \left\| \mathbf{w}^{(t+1)} - \mathbf{w}^* \right\|^2 \right) + \frac{\eta}{2} \|\mathbf{v}_t\|^2 \quad (12)$$

$$\sum_{t=1}^T \langle \mathbf{w}^{(t)} - \mathbf{w}^*, \mathbf{v}_t \rangle = \sum_{t=1}^T \frac{1}{2\eta} \left( \left\| \mathbf{w}^{(t)} - \mathbf{w}^* \right\|^2 - \left\| \mathbf{w}^{(t+1)} - \mathbf{w}^* \right\|^2 \right) + \frac{\eta}{2} \sum_{t=1}^T \|\mathbf{v}_t\|^2 \quad (13)$$

we can find:

$$\sum_{t=1}^T \frac{1}{2\eta} \left( \left\| \mathbf{w}^{(t)} - \mathbf{w}^* \right\|^2 - \left\| \mathbf{w}^{(t+1)} - \mathbf{w}^* \right\|^2 \right) = \frac{1}{2\eta} \left( \left\| \mathbf{w}^{(1)} - \mathbf{w}^* \right\|^2 - \left\| \mathbf{w}^{(T+1)} - \mathbf{w}^* \right\|^2 \right) \leq \frac{\|\mathbf{w}^*\|^2}{2\eta} \quad (14)$$

Thus we can prove:

$$\sum_{t=1}^T \langle \mathbf{w}^{(t)} - \mathbf{w}^*, \mathbf{v}_t \rangle \leq \frac{\|\mathbf{w}^*\|^2}{2\eta} + \frac{\eta}{2} \sum_{t=1}^T \|\mathbf{v}_t\|^2 \quad (15)$$

3. (3 points) Now let's start putting things together and analyze the convergence rate of gradient descent *i.e.* how fast it converges to  $\mathbf{w}^*$ .

First, show that for  $\bar{\mathbf{w}} = \frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)}$

$$f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) \leq \frac{1}{T} \sum_{t=1}^T \langle \mathbf{w}^{(t)} - \mathbf{w}^*, \nabla f(\mathbf{w}^{(t)}) \rangle \quad (16)$$

Next, use the result from part 2, with upper bounds  $B$  and  $\rho$  for  $\|\mathbf{w}^*\|$  and  $\|\nabla f(\mathbf{w}^{(t)})\|$  respectively and show that for fixed  $\eta = \sqrt{\frac{B^2}{\rho^2 T}}$ , the convergence rate of gradient descent is  $\mathcal{O}(1/\sqrt{T})$  *i.e.* the upper bound for  $f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) \propto \frac{1}{\sqrt{T}}$ .

Answer:

$$f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) = f\left(\frac{1}{T} \sum_{t=1}^T \mathbf{w}^{(t)}\right) - f(\mathbf{w}^*) \quad (17)$$

$$\leq \frac{1}{T} f\left(\sum_{t=1}^T \mathbf{w}^{(t)}\right) - f(\mathbf{w}^*) \quad (18)$$

Because of the convexity in question 1, we have:

$$f\left(\sum_{t=1}^T \mathbf{w}^{(t)}\right) - f(\mathbf{w}^*) \leq \langle \sum_{t=1}^T \mathbf{w}^{(t)} - T\mathbf{w}^*, \nabla f(\mathbf{w}^{(t)}) \rangle \quad (19)$$

Thus we can prove:

$$f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) \leq \frac{1}{T} \sum_{t=1}^T \langle \mathbf{w}^{(t)} - \mathbf{w}^*, \nabla f(\mathbf{w}^{(t)}) \rangle \quad (20)$$

Using the conclusion from part2, we can get:

$$f(\bar{\mathbf{w}}) - f(\mathbf{w}^*) \leq \frac{1}{T} \left( \frac{\|\mathbf{w}^*\|^2}{2\eta} + \frac{\eta}{2} \sum_{t=1}^T \left\| \nabla f(\mathbf{w}^{(t)}) \right\|^2 \right) \quad (21)$$

$$\leq \frac{1}{T} \left( \frac{B^2}{2\eta} + \frac{\eta}{2} Tp \right) \quad (22)$$

$$= \frac{B^2}{2\eta T} + \frac{\eta p}{2} \quad (23)$$

$$= \frac{1}{\sqrt{T}} \frac{BP + P}{2} \quad (24)$$

Thus the convergence rate of gradient descent is  $\mathcal{O}(1/\sqrt{T})$

4. (2 points) Consider an objective function  $f(w) := f_1(w) + f_2(w)$  comprised of  $N = 2$  terms:

$$f_1(w) = -\ln \left( 1 - \frac{1}{1 + \exp(-w)} \right) \quad \text{and} \quad f_2(w) = -\ln \left( \frac{1}{1 + \exp(-w)} \right) \quad (25)$$

Now consider using SGD (with a batch-size  $B = 1$ ) to minimize  $f(w)$ . Specifically, in each iteration, we will pick one of the two terms (uniformly at random), and take a step in the direction of the negative gradient, with a constant step-size of  $\eta$ . You can assume  $\eta$  is small enough that every update does result in improvement (aka descent) on the sampled term. Is SGD guaranteed to decrease the overall loss function in every iteration? If yes, provide a proof. If no, provide a counter-example.

Answer: No, SGD does not guarantee to decrease the objective function at every iteration. For example, let  $w^{(0)} = 0$ ,  $w^{(1)} = w^{(0)} - \eta \left( 1 - \frac{1}{(2e^{-w_0}(1+e^{-w_0})^3)} \right) = -\frac{\eta}{2}$ . Since  $\eta$  is positive,  $f(w^{(1)}) > f(w^{(0)})$

## 2 Automatic Differentiation

5. (4 points) In practice, writing the closed-form expression of the derivative of a loss function  $f$  w.r.t. the parameters of a deep neural network is hard (and mostly unnecessary) as  $f$  becomes complex. Instead, we define computation graphs and use the automatic differentiation algorithms (typically backpropagation) to compute gradients using the chain rule. For example, consider the expression

$$f(x, y) = (x + y)(y + 1) \quad (26)$$

Let's define intermediate variables  $a$  and  $b$  such that

$$a = x + y \quad (27)$$

$$b = y + 1 \quad (28)$$

$$f = a \times b \quad (29)$$

A computation graph for the “forward pass” through  $f$  is shown in Fig. 1.




Figure 1

We can then work backwards and compute the derivative of  $f$  w.r.t. each intermediate variable ( $\frac{\partial f}{\partial a}$  and  $\frac{\partial f}{\partial b}$ ) and chain them together to get  $\frac{\partial f}{\partial x}$  and  $\frac{\partial f}{\partial y}$ .

Let  $\sigma(\cdot)$  denote the standard sigmoid function. Now, for the following vector function:

$$f_1(w_1, w_2) = e^{e^{w_1} + e^{2w_2}} + \sigma(e^{w_1} + e^{2w_2}) \quad (30)$$

$$f_2(w_1, w_2) = w_1 w_2 + \max(w_1, w_2) \quad (31)$$

- (a) Draw the computation graph. Compute the value of  $f$  at  $\vec{w} = (1, -1)$ .




Figure 2

- (b) At this  $\vec{w}$ , compute the Jacobian  $\frac{\partial \vec{f}}{\partial \vec{w}}$  using numerical differentiation (using  $\Delta w = 0.01$ ).

Answer:

$$\begin{bmatrix} \frac{\partial f_1}{\partial w_1} & \frac{\partial f_1}{\partial w_2} \\ \frac{\partial f_2}{\partial w_1} & \frac{\partial f_2}{\partial w_2} \end{bmatrix} = \begin{bmatrix} \frac{f_1(w_1+0.01, w_2) - f_1(w_1, w_2)}{0.01} & \frac{f_2(w_1, w_2+0.01) - f_2(w_1, w_2)}{0.01} \\ \frac{f_2(w_1+0.01, w_2) - f_2(w_1, w_2)}{0.01} & \frac{f_2(w_1, w_2+0.01) - f_2(w_1, w_2)}{0.01} \end{bmatrix} = \begin{bmatrix} 48.19 & 4.76 \\ 0 & 1 \end{bmatrix}$$

- (c) At this  $\vec{w}$ , compute the Jacobian using forward mode auto-differentiation.

Answer:

$$y_1 = e^{w_1} \quad (32)$$

$$y_2 = e^{w_1} + e^{2w_2} \quad (33)$$

$$y_3 = e^{y_2} \quad (34)$$

$$y_4 = \sigma(y_3) \quad (35)$$

$$\frac{\partial f_1}{\partial w_1} = \frac{\partial y_5}{\partial w_1} \quad (36)$$

$$= \frac{\partial y_3}{\partial w_1} + \frac{\partial y_4}{\partial w_1} \quad (37)$$

$$= 47.28739 \quad (38)$$

Similarly,

$$\frac{\partial f_1}{\partial w_2} = 4.710 \quad (39)$$

$$\frac{\partial f_2}{\partial w_1} = 1 \quad (40)$$

$$\frac{\partial f_2}{\partial w_2} = 0 \quad (41)$$

$$\begin{bmatrix} 47.28739 & 4.71 \\ 0 & 1 \end{bmatrix}$$

- (d) At this  $\vec{w}$ , compute the Jacobian using backward mode auto-differentiation.

Answer:

$$\frac{\partial f_1}{\partial w_1} = \vec{a}_1 \vec{b}_1 \quad (42)$$

$$= \vec{a}_1 \frac{\exp(a)}{1 + \exp(a)^2} + \vec{d}_2 \exp(c_2) \quad (43)$$

$$= 47.28739 \quad (44)$$

Similarly,

$$\begin{bmatrix} 47.28739 & 4.71 \\ 0 & 1 \end{bmatrix}$$

- (e) Don't you love that software exists to do this for us?

### 3 Paper Review

The first of our paper reviews for this course comes from a much acclaimed spotlight presentation at NeurIPS 2019 on the topic ‘Weight Agnostic Neural Networks’ by Adam Gaier and David Ha from Google Brain.

The paper presents a very interesting proposition that, through a series of experiments, re-examines some fundamental notions about neural networks - in particular, the comparative importance of architectures and weights in a network’s predictive performance.

The paper can be viewed [here](#). The authors have also written a [blog post](#) with intuitive visualizations to help understand its key concepts better.

**Guidelines:** Please restrict your reviews to no more than 350 words. The evaluation rubric for this section is as follows :

6. (2 points) What is the main contribution of this paper? Briefly summarize its key insights, strengths and weaknesses.

Contributions:

In this paper, author starts using deemphasizing weights to search neural network. They aim to search for weight agnostic neural networks, architectures with strong inductive biases that can already perform various tasks with random weights.

First they created an initial population having minimal neural network topologies. Each rollout has shared weight values and can be used to evaluate the network. Then they can rank these networks. They repeat evaluation and eventually they can create a new population.

They use 3 models to test, which is CartPoleSwingUp, BipedalWalker-v2 and CarRacing-v0. Researchers found the results were surprisingly good, as the WANN models with the best-performing shared weight values reached an upright pole position on the CartPoleSwingUp task after only after a few swings. Experiment results also proved that WANNs are no match for convolutional neural networks, which was an expected outcome.

7. (2 points) What is your personal takeaway from this paper? This could be expressed either in terms of relating the approaches adopted in this paper to your traditional understanding of learning parameterized models, or potential future directions of research in the area which the authors haven’t addressed, or anything else that struck you as being noteworthy.

Personal takeaway:

This is a brand new method for searching neural network without using gradient descent. It’s not like traditional neural network and it may lessen the computational resources. As with the age-old nature versus nurture debate, AI researchers want to know whether architecture or weights play the main role in the performance of neural networks. This paper definitely provide a promising start. For me, I think it is interesting but still need to bring it to actual practice to show its performance of the untrained neural network.

### 4 Implement and train a network on CIFAR-10

**Setup Instructions:** Before attempting this question, look at setup instructions at [here](#).

8. (Upto 29 points) Now, we will learn how to implement a softmax classifier, vanilla neural networks (or Multi-Layer Perceptrons), and ConvNets. You will begin by writing the forward and backward passes for different types of layers (including convolution and pooling), and

then go on to train a shallow ConvNet on the CIFAR-10 dataset in Python. Next you will learn to use PyTorch, a popular open-source deep learning framework, and use it to replicate the experiments from before.

Follow the instructions provided [here](#)



# softmax

February 10, 2020

## 1 Softmax Classifier

This exercise guides you through the process of classifying images using a Softmax classifier. As part of this you will:

- Implement a fully vectorized loss function for the Softmax classifier
- Calculate the analytical gradient using vectorized code
- Tune hyperparameters on a validation set
- Optimize the loss function with Stochastic Gradient Descent (SGD)
- Visualize the learned weights

```
[127]: # start-up code!
import random

import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading extenrnal modules
# see http://stackoverflow.com/questions/1907993/
  ↳ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```
[128]: from load_cifar10_tvt import load_cifar10_train_val

X_train, y_train, X_val, y_val, X_test, y_test = load_cifar10_train_val()
print("Train data shape: ", X_train.shape)
print("Train labels shape: ", y_train.shape)
print("Val data shape: ", X_val.shape)
print("Val labels shape: ", y_val.shape)
```

```
print("Test data shape: ", X_test.shape)
print("Test labels shape: ", y_test.shape)
```

Train, validation and testing sets have been created as

$X_i$  and  $y_i$  where  $i=\text{train, val, test}$

Train data shape: (3073, 49000)

Train labels shape: (49000,)

Val data shape: (3073, 1000)

Val labels shape: (1000,)

Test data shape: (3073, 1000)

Test labels shape: (1000,)

Code for this section is to be written in `cs231n/classifiers/softmax.py`

```
[144]: # Now, implement the vectorized version in softmax_loss_vectorized.

import time

from cs231n.classifiers.softmax import softmax_loss_vectorized

# gradient check.
from cs231n.gradient_check import grad_check_sparse

W = np.random.randn(10, 3073) * 0.0001

tic = time.time()
loss, grad = softmax_loss_vectorized(W, X_train, y_train, 0.00001)
toc = time.time()
print("vectorized loss: %e computed in %fs" % (loss, toc - tic))

# As a rough sanity check, our loss should be something close to -log(0.1).
print("loss: %f" % loss)
print("sanity check: %f" % (-np.log(0.1)))

f = lambda w: softmax_loss_vectorized(w, X_train, y_train, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

vectorized loss: 2.383807e+00 computed in 0.147605s

loss: 2.383807

sanity check: 2.302585

numerical: -2.929964 analytic: -2.929964, relative error: 4.616225e-09

numerical: -3.135085 analytic: -3.135085, relative error: 2.552106e-08

numerical: -0.570058 analytic: -0.570058, relative error: 1.404715e-07

numerical: 0.337006 analytic: 0.337006, relative error: 4.889899e-08

numerical: 1.682049 analytic: 1.682049, relative error: 2.496364e-08

numerical: -0.297519 analytic: -0.297519, relative error: 8.559922e-08

numerical: 0.558544 analytic: 0.558544, relative error: 2.837774e-08


numerical: 0.629610 analytic: 0.629610, relative error: 2.700313e-08

numerical: -2.363835 analytic: -2.363835, relative error: 1.094664e-08  
numerical: -1.217689 analytic: -1.217689, relative error: 2.245683e-09

Code for this section is to be written in `cs231n/classifiers/linear_classifier.py`

```
[219]: # Now that efficient implementations to calculate loss function and gradient of  
↪ the softmax are ready,  
# use it to train the classifier on the cifar-10 data  
# Complete the `train` function in cs231n/classifiers/linear_classifier.py  
  
from cs231n.classifiers.linear_classifier import Softmax  
  
classifier = Softmax()  
loss_hist = classifier.train(  
    X_train,  
    y_train,  
    learning_rate=1e-6,  
    reg=1e-5,  
    num_iters=100,  
    batch_size=200,  
    verbose=False,  
)  
# Plot loss vs. iterations  
plt.plot(loss_hist)  
plt.xlabel("Iteration number")  
plt.ylabel("Loss value")
```

```
[219]: Text(0, 0.5, 'Loss value')
```



```
[220]: # Complete the `predict` function in cs231n/classifiers/linear_classifier.py
# Evaluate on test set
y_test_pred = classifier.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print("softmax on raw pixels final test set accuracy: %f" % (test_accuracy,))
```

softmax on raw pixels final test set accuracy: 0.379000

```
[221]: # Visualize the learned weights for each class
w = classifier.W[:, :-1] # strip out the bias
w = w.reshape(10, 32, 32, 3)

w_min, w_max = np.min(w), np.max(w)

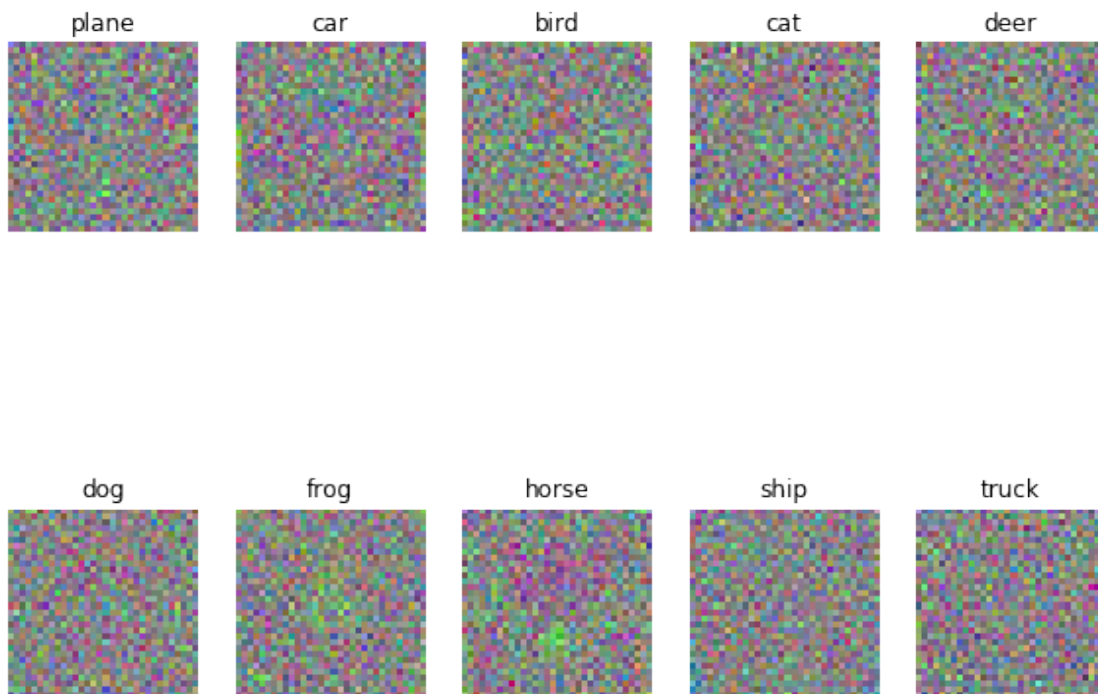
classes = [
    "plane",
    "car",
    "bird",
    "cat",
```

```

    "deer",
    "dog",
    "frog",
    "horse",
    "ship",
    "truck",
]
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype("uint8"))
    plt.axis("off")
    plt.title(classes[i])

```



[ ]:

# two\_layer\_net

February 10, 2020

## 1 Implementing a Neural Network

In this exercise we will develop a neural network with fully-connected layers to perform classification, and test it out on the CIFAR-10 dataset.

```
[1]: # A bit of setup

import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# → autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The neural network parameters will be stored in a dictionary (`model` below), where the keys are the parameter names and the values are numpy arrays. Below, we initialize toy data and a toy model that we will use to verify your implementations.

```
[2]: # Create some toy data to check your implementations

input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5

def init_toy_model():
    model = {}
```

```

    model['W1'] = np.linspace(-0.2, 0.6, num=input_size*hidden_size).
    ↪reshape(input_size, hidden_size)
    model['b1'] = np.linspace(-0.3, 0.7, num=hidden_size)
    model['W2'] = np.linspace(-0.4, 0.1, num=hidden_size*num_classes).
    ↪reshape(hidden_size, num_classes)
    model['b2'] = np.linspace(-0.5, 0.9, num=num_classes)
    return model

def init_toy_data():
    X = np.linspace(-0.2, 0.5, num=num_inputs*input_size).reshape(num_inputs,
    ↪input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

model = init_toy_model()
X, y = init_toy_data()

```

## 2 Forward pass: compute scores

Open the file `cs231n/classifiers/neural_net.py` and look at the function `two_layer_net`. This function is very similar to the loss functions you have written for the Softmax exercise in HW0: It takes the data and weights and computes the class scores, the loss, and the gradients on the parameters.

Implement the first part of the forward pass which uses the weights and biases to compute the scores for all inputs.

```

[12]: from cs231n.classifiers.neural_net import two_layer_net

scores = two_layer_net(X, model)
print(scores)
correct_scores = [[-0.5328368, 0.20031504, 0.93346689],
                  [-0.59412164, 0.15498488, 0.9040914 ],
                  [-0.67658362, 0.08978957, 0.85616275],
                  [-0.77092643, 0.01339997, 0.79772637],
                  [-0.89110401, -0.08754544, 0.71601312]]

# the difference should be very small. We get 3e-8
print('Difference between your scores and correct scores:')
print(np.sum(np.abs(scores - correct_scores)))

```

```

[[-0.5328368  0.20031504  0.93346689]
 [-0.59412164  0.15498488  0.9040914 ]
 [-0.67658362  0.08978957  0.85616275]
 [-0.77092643  0.01339997  0.79772637]
 [-0.89110401 -0.08754544  0.71601312]]

```

Difference between your scores and correct scores:  
3.848682278081994e-08

### 3 Forward pass: compute loss

In the same function, implement the second part that computes the data and regularization loss.

```
[13]: reg = 0.1
      loss, _ = two_layer_net(X, model, y, reg)
      correct_loss = 1.38191946092

      # should be very small, we get 5e-12
      print('Difference between your loss and correct loss:')
      print(np.sum(np.abs(loss - correct_loss)))
```

Difference between your loss and correct loss:  
4.6769255135359344e-12

### 4 Backward pass

Implement the rest of the function. This will compute the gradient of the loss with respect to the variables W1, b1, W2, and b2. Now that you (hopefully!) have a correctly implemented forward pass, you can debug your backward pass using a numeric gradient check:

```
[14]: from cs231n.gradient_check import eval_numerical_gradient

      # Use numeric gradient checking to check your implementation of the backward
      # pass.
      # If your implementation is correct, the difference between the numeric and
      # analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.

      loss, grads = two_layer_net(X, model, y, reg)

      # these should all be less than 1e-8 or so
      for param_name in grads:
          param_grad_num = eval_numerical_gradient(lambda W: two_layer_net(X, model, y,
          reg)[0], model[param_name], verbose=False)
          print('%s max relative error: %e' % (param_name, rel_error(param_grad_num,
          grads[param_name])))
```

W2 max relative error: 9.913918e-10  
b2 max relative error: 8.190173e-11  
W1 max relative error: 4.426512e-09  
b1 max relative error: 5.435431e-08



## 5 Train the network

To train the network we will use SGD with Momentum. Last assignment you implemented vanilla SGD. You will now implement the momentum update and the RMSProp update. Open the file `classifier_trainer.py` and familiarize yourself with the `ClassifierTrainer` class. It performs optimization given an arbitrary cost function data, and model. By default it uses vanilla SGD, which we have already implemented for you. First, run the optimization below using Vanilla SGD:

```
[17]: from cs231n.classifier_trainer import ClassifierTrainer

model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient
↪Descent (no sampled batches of data)
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                              model, two_layer_net,
                                              reg=0.001,
                                              learning_rate=1e-1, momentum=0.0,
↪learning_rate_decay=1,
                                              update='sgd', sample_batches=False,
                                              num_epochs=100,
                                              verbose=False)
print('Final loss with vanilla SGD: %f' % (loss_history[-1], ))

starting iteration  0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with vanilla SGD: 0.940686
```

Now fill in the **momentum update** in the first missing code block inside the `train` function, and run the same optimization as above but with the momentum update. You should see a much better result in the final obtained loss:

```
[18]: model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient
↪Descent (no sampled batches of data)
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                              model, two_layer_net,
```

```

reg=0.001,
learning_rate=1e-1, momentum=0.9,
↪learning_rate_decay=1,
↪sample_batches=False,

update='momentum',

num_epochs=100,
verbose=False)

correct_loss = 0.494394
print('Final loss with momentum SGD: %f. We get: %f' % (loss_history[-1],
↪correct_loss))

```

```

starting iteration 0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with momentum SGD: 0.494394. We get: 0.494394

```

The **RMSProp** update step is given as follows:

```

cache = decay_rate * cache + (1 - decay_rate) * dx**2
x += - learning_rate * dx / np.sqrt(cache + 1e-8)

```

Here, `decay_rate` is a hyperparameter and typical values are [0.9, 0.99, 0.999].

Implement the **RMSProp** update rule inside the `train` function and rerun the optimization:

```

[19]: model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient
↪Descent (no sampled batches of data)
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
reg=0.001,
learning_rate=1e-1, momentum=0.9,
↪learning_rate_decay=1,
↪sample_batches=False,

update='rmsprop',

num_epochs=100,
verbose=False)

correct_loss = 0.439368
print('Final loss with RMSProp: %f. We get: %f' % (loss_history[-1],
↪correct_loss))

```

```
starting iteration 0
starting iteration 10
starting iteration 20
starting iteration 30
starting iteration 40
starting iteration 50
starting iteration 60
starting iteration 70
starting iteration 80
starting iteration 90
Final loss with RMSProp: 0.439368. We get: 0.439368
```

## 6 Load the data

Now that you have implemented a two-layer network that passes gradient checks, it's time to load up our favorite CIFAR-10 data so we can use it to train a classifier.

```
[20]: from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = range(num_training, num_training + num_validation)
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = range(num_training)
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = range(num_test)
    X_test = X_test[mask]
    y_test = y_test[mask]

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis=0)
    X_train -= mean_image
    X_val -= mean_image
    X_test -= mean_image

    # Reshape data to rows
```

```

X_train = X_train.reshape(num_training, -1)
X_val = X_val.reshape(num_validation, -1)
X_test = X_test.reshape(num_test, -1)

return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

```

```

Train data shape: (49000, 3072)
Train labels shape: (49000,)
Validation data shape: (1000, 3072)
Validation labels shape: (1000,)
Test data shape: (1000, 3072)
Test labels shape: (1000,)

```

## 7 Train a network

To train our network we will use SGD with momentum. In addition, we will adjust the learning rate with an exponential learning rate schedule as optimization proceeds; after each epoch, we will reduce the learning rate by multiplying it by a decay rate.

```

[21]: from cs231n.classifiers.neural_net import init_two_layer_model

model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number
↳ of classes
trainer = ClassifierTrainer()
best_model, loss_history, train_acc, val_acc = trainer.train(X_train, y_train,
↳ X_val, y_val,

                                model, two_layer_net,
                                num_epochs=5, reg=1.0,
                                momentum=0.9, learning_rate_decay
↳ = 0.95,

                                learning_rate=1e-5, verbose=True)

```

```

starting iteration 0
Finished epoch 0 / 5: cost 2.302593, train: 0.094000, val 0.108000, lr
1.000000e-05
starting iteration 10

```

starting iteration 20  
starting iteration 30  
starting iteration 40  
starting iteration 50  
starting iteration 60  
starting iteration 70  
starting iteration 80  
starting iteration 90  
starting iteration 100  
starting iteration 110  
starting iteration 120  
starting iteration 130  
starting iteration 140  
starting iteration 150  
starting iteration 160  
starting iteration 170  
starting iteration 180  
starting iteration 190  
starting iteration 200  
starting iteration 210  
starting iteration 220  
starting iteration 230  
starting iteration 240  
starting iteration 250  
starting iteration 260  
starting iteration 270  
starting iteration 280  
starting iteration 290  
starting iteration 300  
starting iteration 310  
starting iteration 320  
starting iteration 330  
starting iteration 340  
starting iteration 350  
starting iteration 360  
starting iteration 370  
starting iteration 380  
starting iteration 390  
starting iteration 400  
starting iteration 410  
starting iteration 420  
starting iteration 430  
starting iteration 440  
starting iteration 450  
starting iteration 460  
starting iteration 470  
starting iteration 480  
Finished epoch 1 / 5: cost 2.285003, train: 0.193000, val 0.173000, lr

9.500000e-06

starting iteration 490  
starting iteration 500  
starting iteration 510  
starting iteration 520  
starting iteration 530  
starting iteration 540  
starting iteration 550  
starting iteration 560  
starting iteration 570  
starting iteration 580  
starting iteration 590  
starting iteration 600  
starting iteration 610  
starting iteration 620  
starting iteration 630  
starting iteration 640  
starting iteration 650  
starting iteration 660  
starting iteration 670  
starting iteration 680  
starting iteration 690  
starting iteration 700  
starting iteration 710  
starting iteration 720  
starting iteration 730  
starting iteration 740  
starting iteration 750  
starting iteration 760  
starting iteration 770  
starting iteration 780  
starting iteration 790  
starting iteration 800  
starting iteration 810  
starting iteration 820  
starting iteration 830  
starting iteration 840  
starting iteration 850  
starting iteration 860  
starting iteration 870  
starting iteration 880  
starting iteration 890  
starting iteration 900  
starting iteration 910  
starting iteration 920  
starting iteration 930  
starting iteration 940  
starting iteration 950

starting iteration 960  
starting iteration 970  
Finished epoch 2 / 5: cost 2.154377, train: 0.256000, val 0.240000, lr  
9.025000e-06  
starting iteration 980  
starting iteration 990  
starting iteration 1000  
starting iteration 1010  
starting iteration 1020  
starting iteration 1030  
starting iteration 1040  
starting iteration 1050  
starting iteration 1060  
starting iteration 1070  
starting iteration 1080  
starting iteration 1090  
starting iteration 1100  
starting iteration 1110  
starting iteration 1120  
starting iteration 1130  
starting iteration 1140  
starting iteration 1150  
starting iteration 1160  
starting iteration 1170  
starting iteration 1180  
starting iteration 1190  
starting iteration 1200  
starting iteration 1210  
starting iteration 1220  
starting iteration 1230  
starting iteration 1240  
starting iteration 1250  
starting iteration 1260  
starting iteration 1270  
starting iteration 1280  
starting iteration 1290  
starting iteration 1300  
starting iteration 1310  
starting iteration 1320  
starting iteration 1330  
starting iteration 1340  
starting iteration 1350  
starting iteration 1360  
starting iteration 1370  
starting iteration 1380  
starting iteration 1390  
starting iteration 1400  
starting iteration 1410

starting iteration 1420  
starting iteration 1430  
starting iteration 1440  
starting iteration 1450  
starting iteration 1460  
Finished epoch 3 / 5: cost 1.928161, train: 0.264000, val 0.282000, lr  
8.573750e-06  
starting iteration 1470  
starting iteration 1480  
starting iteration 1490  
starting iteration 1500  
starting iteration 1510  
starting iteration 1520  
starting iteration 1530  
starting iteration 1540  
starting iteration 1550  
starting iteration 1560  
starting iteration 1570  
starting iteration 1580  
starting iteration 1590  
starting iteration 1600  
starting iteration 1610  
starting iteration 1620  
starting iteration 1630  
starting iteration 1640  
starting iteration 1650  
starting iteration 1660  
starting iteration 1670  
starting iteration 1680  
starting iteration 1690  
starting iteration 1700  
starting iteration 1710  
starting iteration 1720  
starting iteration 1730  
starting iteration 1740  
starting iteration 1750  
starting iteration 1760  
starting iteration 1770  
starting iteration 1780  
starting iteration 1790  
starting iteration 1800  
starting iteration 1810  
starting iteration 1820  
starting iteration 1830  
starting iteration 1840  
starting iteration 1850  
starting iteration 1860  
starting iteration 1870



starting iteration 1880  
starting iteration 1890  
starting iteration 1900  
starting iteration 1910  
starting iteration 1920  
starting iteration 1930  
starting iteration 1940  
starting iteration 1950  
Finished epoch 4 / 5: cost 1.799440, train: 0.329000, val 0.331000, lr  
8.145063e-06  
starting iteration 1960  
starting iteration 1970  
starting iteration 1980  
starting iteration 1990  
starting iteration 2000  
starting iteration 2010  
starting iteration 2020  
starting iteration 2030  
starting iteration 2040  
starting iteration 2050  
starting iteration 2060  
starting iteration 2070  
starting iteration 2080  
starting iteration 2090  
starting iteration 2100  
starting iteration 2110  
starting iteration 2120  
starting iteration 2130  
starting iteration 2140  
starting iteration 2150  
starting iteration 2160  
starting iteration 2170  
starting iteration 2180  
starting iteration 2190  
starting iteration 2200  
starting iteration 2210  
starting iteration 2220  
starting iteration 2230  
starting iteration 2240  
starting iteration 2250  
starting iteration 2260  
starting iteration 2270  
starting iteration 2280  
starting iteration 2290  
starting iteration 2300  
starting iteration 2310  
starting iteration 2320  
starting iteration 2330

```
starting iteration 2340
starting iteration 2350
starting iteration 2360
starting iteration 2370
starting iteration 2380
starting iteration 2390
starting iteration 2400
starting iteration 2410
starting iteration 2420
starting iteration 2430
starting iteration 2440
Finished epoch 5 / 5: cost 1.724871, train: 0.351000, val 0.356000, lr
7.737809e-06
finished optimization. best validation accuracy: 0.356000
```

## 8 Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 0.37 on the validation set. This isn't very good.


One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
[22]: # Plot the loss function and train / validation accuracies
plt.subplot(2, 1, 1)
plt.plot(loss_history)
plt.title('Loss history')
plt.xlabel('Iteration')
plt.ylabel('Loss')

plt.subplot(2, 1, 2)
plt.plot(train_acc)
plt.plot(val_acc)
plt.legend(['Training accuracy', 'Validation accuracy'], loc='lower right')
plt.xlabel('Epoch')
plt.ylabel('Clasification accuracy')
```

```
[22]: Text(0, 0.5, 'Clasification accuracy')
```




```
[23]: from cs231n.vis_utils import visualize_grid

# Visualize the weights of the network

def show_net_weights(model):
    plt.imshow(visualize_grid(model['W1'].T.reshape(-1, 32, 32, 3), padding=3).
        ↪astype('uint8'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(model)
```



## 9 Tune your hyperparameters

**What's wrong?.** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer

size, learning rate, number of training epochs, and regularization strength. You might also consider tuning the momentum and learning rate decay parameters, but you should be able to get good performance using the default values.

**Approximate results.** You should be able to achieve a classification accuracy of greater than 50% on the validation set. Our best network gets over 56% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on CIFAR-10 as you can, with a fully-connected Neural Network. For every 1% above 56% on the Test set we will award you with one extra bonus point. Feel free to implement your own techniques (e.g. PCA to reduce dimensionality, or adding dropout, or adding features to the solver, etc.).

```
[28]: best_model = None # store the best model into this

#####
# TODO: Tune hyperparameters using the validation set. Store your best trained
# model in best_model.
#
# To help debug your network, it may help to use visualizations similar to the
# ones we used above; these visualizations will have significant qualitative
# differences from the ones we saw above for the poorly tuned network.
#
# Tweaking hyperparameters by hand can be fun, but you might find it useful to
# write code to sweep through possible combinations of hyperparameters
# automatically like we did on the previous assignment.
#####
# input size, hidden size, number of classes
model = init_two_layer_model(32*32*3, 1000, 10)
trainer = ClassifierTrainer()
best_model, loss_history, train_acc, val_acc = trainer.train(X_train, y_train,
                                                             X_val, y_val,
                                                             model, two_layer_net,
                                                             num_epochs=10, reg=0.0001,
                                                             momentum=0.9,
                                                             learning_rate_decay=0.7,
                                                             learning_rate=1e-4, verbose=True)
#####
```

```
#
```

```
END OF YOUR CODE
```

```
␣
```

```
↪#
```

```
#####
```

```
starting iteration 0
```

```
Finished epoch 0 / 10: cost 2.302586, train: 0.137000, val 0.105000, lr  
1.000000e-04
```

```
starting iteration 10
```

```
starting iteration 20
```

```
starting iteration 30
```

```
starting iteration 40
```

```
starting iteration 50
```

```
starting iteration 60
```

```
starting iteration 70
```

```
starting iteration 80
```

```
starting iteration 90
```

```
starting iteration 100
```

```
starting iteration 110
```

```
starting iteration 120
```

```
starting iteration 130
```

```
starting iteration 140
```

```
starting iteration 150
```

```
starting iteration 160
```

```
starting iteration 170
```

```
starting iteration 180
```

```
starting iteration 190
```

```
starting iteration 200
```

```
starting iteration 210
```

```
starting iteration 220
```

```
starting iteration 230
```

```
starting iteration 240
```

```
starting iteration 250
```

```
starting iteration 260
```

```
starting iteration 270
```

```
starting iteration 280
```

```
starting iteration 290
```

```
starting iteration 300
```

```
starting iteration 310
```

```
starting iteration 320
```

```
starting iteration 330
```

```
starting iteration 340
```

```
starting iteration 350
```

```
starting iteration 360
```

```
starting iteration 370
```

```
starting iteration 380
```

```
starting iteration 390
```

```
starting iteration 400
```

starting iteration 410  
starting iteration 420  
starting iteration 430  
starting iteration 440  
starting iteration 450  
starting iteration 460  
starting iteration 470  
starting iteration 480  
Finished epoch 1 / 10: cost 1.498791, train: 0.451000, val 0.450000, lr  
7.000000e-05  
starting iteration 490  
starting iteration 500  
starting iteration 510  
starting iteration 520  
starting iteration 530  
starting iteration 540  
starting iteration 550  
starting iteration 560  
starting iteration 570  
starting iteration 580  
starting iteration 590  
starting iteration 600  
starting iteration 610  
starting iteration 620  
starting iteration 630  
starting iteration 640  
starting iteration 650  
starting iteration 660  
starting iteration 670  
starting iteration 680  
starting iteration 690  
starting iteration 700  
starting iteration 710  
starting iteration 720  
starting iteration 730  
starting iteration 740  
starting iteration 750  
starting iteration 760  
starting iteration 770  
starting iteration 780  
starting iteration 790  
starting iteration 800  
starting iteration 810  
starting iteration 820  
starting iteration 830  
starting iteration 840  
starting iteration 850  
starting iteration 860

starting iteration 870  
starting iteration 880  
starting iteration 890  
starting iteration 900  
starting iteration 910  
starting iteration 920  
starting iteration 930  
starting iteration 940  
starting iteration 950  
starting iteration 960  
starting iteration 970  
Finished epoch 2 / 10: cost 1.667420, train: 0.510000, val 0.482000, lr  
4.900000e-05  
starting iteration 980  
starting iteration 990  
starting iteration 1000  
starting iteration 1010  
starting iteration 1020  
starting iteration 1030  
starting iteration 1040  
starting iteration 1050  
starting iteration 1060  
starting iteration 1070  
starting iteration 1080  
starting iteration 1090  
starting iteration 1100  
starting iteration 1110  
starting iteration 1120  
starting iteration 1130  
starting iteration 1140  
starting iteration 1150  
starting iteration 1160  
starting iteration 1170  
starting iteration 1180  
starting iteration 1190  
starting iteration 1200  
starting iteration 1210  
starting iteration 1220  
starting iteration 1230  
starting iteration 1240  
starting iteration 1250  
starting iteration 1260  
starting iteration 1270  
starting iteration 1280  
starting iteration 1290  
starting iteration 1300  
starting iteration 1310  
starting iteration 1320



starting iteration 1330  
starting iteration 1340  
starting iteration 1350  
starting iteration 1360  
starting iteration 1370  
starting iteration 1380  
starting iteration 1390  
starting iteration 1400  
starting iteration 1410  
starting iteration 1420  
starting iteration 1430  
starting iteration 1440  
starting iteration 1450  
starting iteration 1460  
Finished epoch 3 / 10: cost 1.220875, train: 0.513000, val 0.484000, lr  
3.430000e-05  
starting iteration 1470  
starting iteration 1480  
starting iteration 1490  
starting iteration 1500  
starting iteration 1510  
starting iteration 1520  
starting iteration 1530  
starting iteration 1540  
starting iteration 1550  
starting iteration 1560  
starting iteration 1570  
starting iteration 1580  
starting iteration 1590  
starting iteration 1600  
starting iteration 1610  
starting iteration 1620  
starting iteration 1630  
starting iteration 1640  
starting iteration 1650  
starting iteration 1660  
starting iteration 1670  
starting iteration 1680  
starting iteration 1690  
starting iteration 1700  
starting iteration 1710  
starting iteration 1720  
starting iteration 1730  
starting iteration 1740  
starting iteration 1750  
starting iteration 1760  
starting iteration 1770  
starting iteration 1780

starting iteration 1790  
starting iteration 1800  
starting iteration 1810  
starting iteration 1820  
starting iteration 1830  
starting iteration 1840  
starting iteration 1850  
starting iteration 1860  
starting iteration 1870  
starting iteration 1880  
starting iteration 1890  
starting iteration 1900  
starting iteration 1910  
starting iteration 1920  
starting iteration 1930  
starting iteration 1940  
starting iteration 1950  
Finished epoch 4 / 10: cost 1.415856, train: 0.548000, val 0.508000, lr  
2.401000e-05  
starting iteration 1960  
starting iteration 1970  
starting iteration 1980  
starting iteration 1990  
starting iteration 2000  
starting iteration 2010  
starting iteration 2020  
starting iteration 2030  
starting iteration 2040  
starting iteration 2050  
starting iteration 2060  
starting iteration 2070  
starting iteration 2080  
starting iteration 2090  
starting iteration 2100  
starting iteration 2110  
starting iteration 2120  
starting iteration 2130  
starting iteration 2140  
starting iteration 2150  
starting iteration 2160  
starting iteration 2170  
starting iteration 2180  
starting iteration 2190  
starting iteration 2200  
starting iteration 2210  
starting iteration 2220  
starting iteration 2230  
starting iteration 2240

starting iteration 2250  
starting iteration 2260  
starting iteration 2270  
starting iteration 2280  
starting iteration 2290  
starting iteration 2300  
starting iteration 2310  
starting iteration 2320  
starting iteration 2330  
starting iteration 2340  
starting iteration 2350  
starting iteration 2360  
starting iteration 2370  
starting iteration 2380  
starting iteration 2390  
starting iteration 2400  
starting iteration 2410  
starting iteration 2420  
starting iteration 2430  
starting iteration 2440  
Finished epoch 5 / 10: cost 1.094622, train: 0.580000, val 0.527000, lr  
1.680700e-05  
starting iteration 2450  
starting iteration 2460  
starting iteration 2470  
starting iteration 2480  
starting iteration 2490  
starting iteration 2500  
starting iteration 2510  
starting iteration 2520  
starting iteration 2530  
starting iteration 2540  
starting iteration 2550  
starting iteration 2560  
starting iteration 2570  
starting iteration 2580  
starting iteration 2590  
starting iteration 2600  
starting iteration 2610  
starting iteration 2620  
starting iteration 2630  
starting iteration 2640  
starting iteration 2650  
starting iteration 2660  
starting iteration 2670  
starting iteration 2680  
starting iteration 2690  
starting iteration 2700

starting iteration 2710  
starting iteration 2720  
starting iteration 2730  
starting iteration 2740  
starting iteration 2750  
starting iteration 2760  
starting iteration 2770  
starting iteration 2780  
starting iteration 2790  
starting iteration 2800  
starting iteration 2810  
starting iteration 2820  
starting iteration 2830  
starting iteration 2840  
starting iteration 2850  
starting iteration 2860  
starting iteration 2870  
starting iteration 2880  
starting iteration 2890  
starting iteration 2900  
starting iteration 2910  
starting iteration 2920  
starting iteration 2930  
Finished epoch 6 / 10: cost 1.398334, train: 0.613000, val 0.532000, lr  
1.176490e-05  
starting iteration 2940  
starting iteration 2950  
starting iteration 2960  
starting iteration 2970  
starting iteration 2980  
starting iteration 2990  
starting iteration 3000  
starting iteration 3010  
starting iteration 3020  
starting iteration 3030  
starting iteration 3040  
starting iteration 3050  
starting iteration 3060  
starting iteration 3070  
starting iteration 3080  
starting iteration 3090  
starting iteration 3100  
starting iteration 3110  
starting iteration 3120  
starting iteration 3130  
starting iteration 3140  
starting iteration 3150  
starting iteration 3160

starting iteration 3170  
starting iteration 3180  
starting iteration 3190  
starting iteration 3200  
starting iteration 3210  
starting iteration 3220  
starting iteration 3230  
starting iteration 3240  
starting iteration 3250  
starting iteration 3260  
starting iteration 3270  
starting iteration 3280  
starting iteration 3290  
starting iteration 3300  
starting iteration 3310  
starting iteration 3320  
starting iteration 3330  
starting iteration 3340  
starting iteration 3350  
starting iteration 3360  
starting iteration 3370  
starting iteration 3380  
starting iteration 3390  
starting iteration 3400  
starting iteration 3410  
starting iteration 3420  
Finished epoch 7 / 10: cost 1.084483, train: 0.562000, val 0.544000, lr  
8.235430e-06  
starting iteration 3430  
starting iteration 3440  
starting iteration 3450  
starting iteration 3460  
starting iteration 3470  
starting iteration 3480  
starting iteration 3490  
starting iteration 3500  
starting iteration 3510  
starting iteration 3520  
starting iteration 3530  
starting iteration 3540  
starting iteration 3550  
starting iteration 3560  
starting iteration 3570  
starting iteration 3580  
starting iteration 3590  
starting iteration 3600  
starting iteration 3610  
starting iteration 3620


starting iteration 3630  
starting iteration 3640  
starting iteration 3650  
starting iteration 3660  
starting iteration 3670  
starting iteration 3680  
starting iteration 3690  
starting iteration 3700  
starting iteration 3710  
starting iteration 3720  
starting iteration 3730  
starting iteration 3740  
starting iteration 3750  
starting iteration 3760  
starting iteration 3770  
starting iteration 3780  
starting iteration 3790  
starting iteration 3800  
starting iteration 3810  
starting iteration 3820  
starting iteration 3830  
starting iteration 3840  
starting iteration 3850  
starting iteration 3860  
starting iteration 3870  
starting iteration 3880  
starting iteration 3890  
starting iteration 3900  
starting iteration 3910  
Finished epoch 8 / 10: cost 1.069327, train: 0.624000, val 0.534000, lr  
5.764801e-06  
starting iteration 3920  
starting iteration 3930  
starting iteration 3940  
starting iteration 3950  
starting iteration 3960  
starting iteration 3970  
starting iteration 3980  
starting iteration 3990  
starting iteration 4000  
starting iteration 4010  
starting iteration 4020  
starting iteration 4030  
starting iteration 4040  
starting iteration 4050  
starting iteration 4060  
starting iteration 4070  
starting iteration 4080

```
starting iteration 4090
starting iteration 4100
starting iteration 4110
starting iteration 4120
starting iteration 4130
starting iteration 4140
starting iteration 4150
starting iteration 4160
starting iteration 4170
starting iteration 4180
starting iteration 4190
starting iteration 4200
starting iteration 4210
starting iteration 4220
starting iteration 4230
starting iteration 4240
starting iteration 4250
starting iteration 4260
starting iteration 4270
starting iteration 4280
starting iteration 4290
starting iteration 4300
starting iteration 4310
starting iteration 4320
starting iteration 4330
starting iteration 4340
starting iteration 4350
starting iteration 4360
starting iteration 4370
starting iteration 4380
starting iteration 4390
starting iteration 4400
Finished epoch 9 / 10: cost 1.126785, train: 0.651000, val 0.557000, lr
4.035361e-06
starting iteration 4410
starting iteration 4420
starting iteration 4430
starting iteration 4440
starting iteration 4450
starting iteration 4460
starting iteration 4470
starting iteration 4480
starting iteration 4490
starting iteration 4500
starting iteration 4510
starting iteration 4520
starting iteration 4530
starting iteration 4540
```

```
starting iteration 4550
starting iteration 4560
starting iteration 4570
starting iteration 4580
starting iteration 4590
starting iteration 4600
starting iteration 4610
starting iteration 4620
starting iteration 4630
starting iteration 4640
starting iteration 4650
starting iteration 4660
starting iteration 4670
starting iteration 4680
starting iteration 4690
starting iteration 4700
starting iteration 4710
starting iteration 4720
starting iteration 4730
starting iteration 4740
starting iteration 4750
starting iteration 4760
starting iteration 4770
starting iteration 4780
starting iteration 4790
starting iteration 4800
starting iteration 4810
starting iteration 4820
starting iteration 4830
starting iteration 4840
starting iteration 4850
starting iteration 4860
starting iteration 4870
starting iteration 4880
starting iteration 4890
Finished epoch 10 / 10: cost 1.272121, train: 0.601000, val 0.543000, lr
2.824752e-06
finished optimization. best validation accuracy: 0.557000
```

```
[29]: # visualize the weights
      show_net_weights(best_model)
```





## 10 Run on the test set

When you are done experimenting, you should evaluate your final trained network on the test set.

```
[30]: scores_test = two_layer_net(X_test, best_model)
      print('Test accuracy: ', np.mean(np.argmax(scores_test, axis=1) == y_test))
```

Test accuracy: 0.538

```
[ ]:
```

# layers

February 10, 2020

## 1 Modular neural nets

In the previous exercise, we computed the loss and gradient for a two-layer neural network in a single monolithic function. This isn't very difficult for a small two-layer network, but would be tedious and error-prone for larger networks. Ideally we want to build networks using a more modular design so that we can snap together different types of layers and loss functions in order to quickly experiment with different architectures.

In this exercise we will implement this approach, and develop a number of different layer types in isolation that can then be easily plugged together. For each layer we will implement `forward` and `backward` functions. The `forward` function will receive data, weights, and other parameters, and will return both an output and a `cache` object that stores data needed for the backward pass. The `backward` function will receive upstream derivatives and the cache object, and will return gradients with respect to the data and all of the weights. This will allow us to write code that looks like this:

```
def two_layer_net(X, W1, b1, W2, b2, reg):
    # Forward pass; compute scores
    s1, fc1_cache = affine_forward(X, W1, b1)
    a1, relu_cache = relu_forward(s1)
    scores, fc2_cache = affine_forward(a1, W2, b2)

    # Loss functions return data loss and gradients on scores
    data_loss, dscores = svm_loss(scores, y)

    # Compute backward pass
    da1, dW2, db2 = affine_backward(dscores, fc2_cache)
    ds1 = relu_backward(da1, relu_cache)
    dX, dW1, db1 = affine_backward(ds1, fc1_cache)

    # A real network would add regularization here

    # Return loss and gradients
    return loss, dW1, db1, dW2, db2
```

```
[1]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
```

```

from cs231n.gradient_check import eval_numerical_gradient_array, \
    eval_numerical_gradient
from cs231n.layers import *

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

```

## 2 Affine layer: forward

Open the file `cs231n/layers.py` and implement the `affine_forward` function.

Once you are done we will test your can test your implementation by running the following:

```

[2]: # Test the affine_forward function

num_inputs = 2
input_shape = (4, 5, 6)
output_dim = 3

input_size = num_inputs * np.prod(input_shape)
weight_size = output_dim * np.prod(input_shape)

x = np.linspace(-0.1, 0.5, num=input_size).reshape(num_inputs, *input_shape)
w = np.linspace(-0.2, 0.3, num=weight_size).reshape(np.prod(input_shape), \
    output_dim)
b = np.linspace(-0.3, 0.1, num=output_dim)

out, _ = affine_forward(x, w, b)
correct_out = np.array([[ 1.49834967,  1.70660132,  1.91485297],
                        [ 3.25553199,  3.5141327,  3.77273342]])

# Compare your output with ours. The error should be around 1e-9.
print('Testing affine_forward function:')
print('difference: ', rel_error(out, correct_out))

```

Testing affine\_forward function:  
difference: 9.769847728806635e-10

### 3 Affine layer: backward

Now implement the `affine_backward` function. You can test your implementation using numeric gradient checking.

```
[3]: # Test the affine_backward function

x = np.random.randn(10, 2, 3)
w = np.random.randn(6, 5)
b = np.random.randn(5)
dout = np.random.randn(10, 5)

dx_num = eval_numerical_gradient_array(lambda x: affine_forward(x, w, b)[0], x,
    ↪dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_forward(x, w, b)[0], w,
    ↪dout)
db_num = eval_numerical_gradient_array(lambda b: affine_forward(x, w, b)[0], b,
    ↪dout)

_, cache = affine_forward(x, w, b)
dx, dw, db = affine_backward(dout, cache)

# The error should be less than 1e-10
print('Testing affine_backward function:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

Testing affine\_backward function:  
dx error: 8.382793903902256e-11  
dw error: 1.558761275742988e-10  
db error: 4.6002441594006085e-11

### 4 ReLU layer: forward

Implement the `relu_forward` function and test your implementation by running the following:

```
[4]: # Test the relu_forward function

x = np.linspace(-0.5, 0.5, num=12).reshape(3, 4)

out, _ = relu_forward(x)
```

```

correct_out = np.array([[ 0.,          0.,          0.,          0.,          ],
                        [ 0.,          0.,          0.04545455, 0.13636364,],
                        [ 0.22727273, 0.31818182, 0.40909091, 0.5,          ]])

# Compare your output with ours. The error should be around 1e-8
print('Testing relu_forward function:')
print('difference: ', rel_error(out, correct_out))

```

Testing relu\_forward function:  
 difference: 4.999999798022158e-08

## 5 ReLU layer: backward

Implement the `relu_backward` function and test your implementation using numeric gradient checking:

```

[5]: x = np.random.randn(10, 10)
dout = np.random.randn(*x.shape)

dx_num = eval_numerical_gradient_array(lambda x: relu_forward(x)[0], x, dout)

_, cache = relu_forward(x)
dx = relu_backward(dout, cache)

# The error should be around 1e-12
print('Testing relu_backward function:')
print('dx error: ', rel_error(dx_num, dx))

```

Testing relu\_backward function:  
 dx error: 3.275618483368625e-12

## 6 Loss layers: Softmax and SVM

You implemented these loss functions in the last assignment, so we'll give them to you for free here. It's still a good idea to test them to make sure they work correctly.

```

[6]: num_classes, num_inputs = 10, 50
x = 0.001 * np.random.randn(num_inputs, num_classes)
y = np.random.randint(num_classes, size=num_inputs)

dx_num = eval_numerical_gradient(lambda x: svm_loss(x, y)[0], x, verbose=False)
loss, dx = svm_loss(x, y)

# Test svm_loss function. Loss should be around 9 and dx error should be 1e-9
print('Testing svm_loss:')

```

```

print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))

dx_num = eval_numerical_gradient(lambda x: softmax_loss(x, y)[0], x,
    verbose=False)
loss, dx = softmax_loss(x, y)

# Test softmax_loss function. Loss should be 2.3 and dx error should be 1e-8
print('\nTesting softmax_loss:')
print('loss: ', loss)
print('dx error: ', rel_error(dx_num, dx))

```

Testing svm\_loss:

loss: 9.000597824304572

dx error: 8.182894472887002e-10

Testing softmax\_loss:

loss: 2.3026453494071113

dx error: 7.485321900683551e-09

## 7 Convolution layer: forward naive

We are now ready to implement the forward pass for a convolutional layer. Implement the function `conv_forward_naive` in the file `cs231n/layers.py`.

You don't have to worry too much about efficiency at this point; just write the code in whatever way you find most clear.

You can test your implementation by running the following:

```

[7]: x_shape = (2, 3, 4, 4)
w_shape = (3, 3, 4, 4)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
b = np.linspace(-0.1, 0.2, num=3)

conv_param = {'stride': 2, 'pad': 1}
out, _ = conv_forward_naive(x, w, b, conv_param)
correct_out = np.array([[[[[-0.08759809, -0.10987781],
                           [-0.18387192, -0.2109216 ]],
                          [[ 0.21027089,  0.21661097],
                           [ 0.22847626,  0.23004637]],
                          [[ 0.50813986,  0.54309974],
                           [ 0.64082444,  0.67101435]]],
                         [[[-0.98053589, -1.03143541],
                           [-1.19128892, -1.24695841]],
                          [[ 0.69108355,  0.66880383],

```



```

[ 0.59480972, 0.56776003]],
[[ 2.36270298, 2.36904306],
 [ 2.38090835, 2.38247847]]]]])

# Compare your output to ours; difference should be around 1e-8
print('Testing conv_forward_naive')
print('difference: ', rel_error(out, correct_out))

```

```

Testing conv_forward_naive
difference: 2.2121476417505994e-08

```

## 8 Aside: Image processing via convolutions

As fun way to both check your implementation and gain a better understanding of the type of operation that convolutional layers can perform, we will set up an input containing two images and manually set up filters that perform common image processing operations (grayscale conversion and edge detection). The convolution forward pass will apply these operations to each of the input images. We can then visualize the results as a sanity check.

```

[8]: from scipy.misc import imread, imresize

kitten, puppy = imread('kitten.jpg'), imread('puppy.jpg')
# kitten is wide, and puppy is already square
d = kitten.shape[1] - kitten.shape[0]
kitten_cropped = kitten[:, d//2:-d//2, :]

img_size = 200 # Make this smaller if it runs too slow
x = np.zeros((2, 3, img_size, img_size))
x[0, :, :, :] = imresize(puppy, (img_size, img_size)).transpose((2, 0, 1))
x[1, :, :, :] = imresize(kitten_cropped, (img_size, img_size)).transpose((2, 0, 1))

# Set up a convolutional weights holding 2 filters, each 3x3
w = np.zeros((2, 3, 3, 3))

# The first filter converts the image to grayscale.
# Set up the red, green, and blue channels of the filter.
w[0, 0, :, :] = [[0, 0, 0], [0, 0.3, 0], [0, 0, 0]]
w[0, 1, :, :] = [[0, 0, 0], [0, 0.6, 0], [0, 0, 0]]
w[0, 2, :, :] = [[0, 0, 0], [0, 0.1, 0], [0, 0, 0]]

# Second filter detects horizontal edges in the blue channel.
w[1, 2, :, :] = [[1, 2, 1], [0, 0, 0], [-1, -2, -1]]

# Vector of biases. We don't need any bias for the grayscale
# filter, but for the edge detection filter we want to add 128

```

```

# to each output so that nothing is negative.
b = np.array([0, 128])

# Compute the result of convolving each input in x with each filter in w,
# offsetting by b, and storing the results in out.
out, _ = conv_forward_naive(x, w, b, {'stride': 1, 'pad': 1})

def imshow_noax(img, normalize=True):
    """ Tiny helper to show images as uint8 and remove axis labels """
    if normalize:
        img_max, img_min = np.max(img), np.min(img)
        img = 255.0 * (img - img_min) / (img_max - img_min)
    plt.imshow(img.astype('uint8'))
    plt.gca().axis('off')

# Show the original images and the results of the conv operation
plt.subplot(2, 3, 1)
imshow_noax(puppy, normalize=False)
plt.title('Original image')
plt.subplot(2, 3, 2)
imshow_noax(out[0, 0])
plt.title('Grayscale')
plt.subplot(2, 3, 3)
imshow_noax(out[0, 1])
plt.title('Edges')
plt.subplot(2, 3, 4)
imshow_noax(kitten_cropped, normalize=False)
plt.subplot(2, 3, 5)
imshow_noax(out[1, 0])
plt.subplot(2, 3, 6)
imshow_noax(out[1, 1])
plt.show()

```

C:\Users\82120\.conda\envs\7643\lib\site-packages\ipykernel\_launcher.py:3:

DeprecationWarning: `imread` is deprecated!

`imread` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.

Use ``imageio.imread`` instead.

This is separate from the ipykernel package so we can avoid doing imports until

C:\Users\82120\.conda\envs\7643\lib\site-packages\ipykernel\_launcher.py:10:

DeprecationWarning: `imresize` is deprecated!

`imresize` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.

Use ``skimage.transform.resize`` instead.

# Remove the CWD from sys.path while we load stuff.

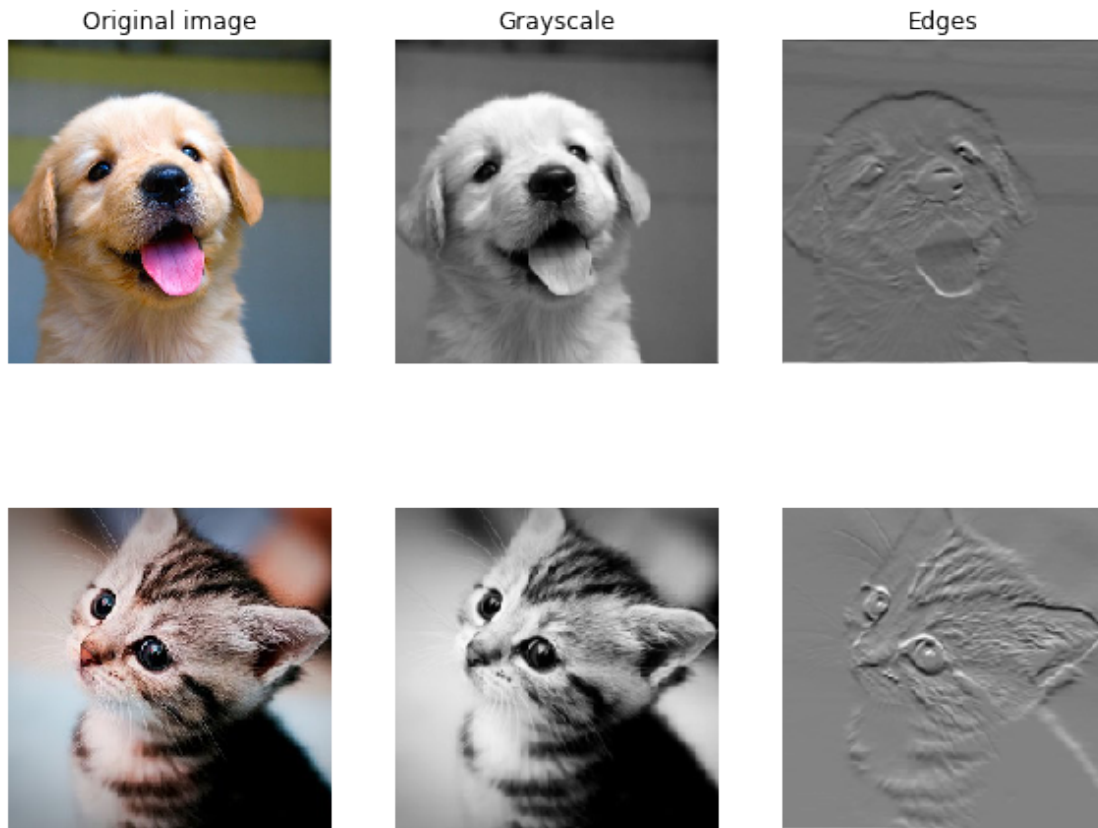
C:\Users\82120\.conda\envs\7643\lib\site-packages\ipykernel\_launcher.py:11:

DeprecationWarning: `imresize` is deprecated!

`imresize` is deprecated in SciPy 1.0.0, and will be removed in 1.2.0.



Use `skimage.transform.resize` instead.  
# This is added back by `InteractiveShellApp.init_path()`



## 9 Convolution layer: backward naive

Next you need to implement the function `conv_backward_naive` in the file `cs231n/layers.py`. As usual, we will check your implementation with numeric gradient checking.

```
[25]: x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
conv_param = {'stride': 1, 'pad': 1}

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w, b,
    ↪conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w, b,
    ↪conv_param)[0], w, dout)
```

```

db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w, b,
↳conv_param)[0], b, dout)

out, cache = conv_forward_naive(x, w, b, conv_param)
dx, dw, db = conv_backward_naive(dout, cache)

# Your errors should be around 1e-9
print('Testing conv_backward_naive function')
print('dx error: ', rel_error(dx, dx_num))
print('dw error: ', rel_error(dw, dw_num))
print('db error: ', rel_error(db, db_num))

```

5

```

Testing conv_backward_naive function
dx error:  3.16025486039339e-09
dw error:  5.03599497021093e-10
db error:  1.2370846328888066e-10

```

## 10 Max pooling layer: forward naive

The last layer we need for a basic convolutional neural network is the max pooling layer. First implement the forward pass in the function `max_pool_forward_naive` in the file `cs231n/layers.py`.

```

[26]: x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)

correct_out = np.array([[[[-0.26315789, -0.24842105],
                           [-0.20421053, -0.18947368]],
                          [[-0.14526316, -0.13052632],
                           [-0.08631579, -0.07157895]],
                          [[-0.02736842, -0.01263158],
                           [ 0.03157895,  0.04631579]]],
                        [[[ 0.09052632,  0.10526316],
                           [ 0.14947368,  0.16421053]],
                          [[ 0.20842105,  0.22315789],
                           [ 0.26736842,  0.28210526]],
                          [[ 0.32631579,  0.34105263],
                           [ 0.38526316,  0.4          ]]]])

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))

```

Testing max\_pool\_forward\_naive function:  
difference: 4.1666665157267834e-08

## 11 Max pooling layer: backward naive

Implement the backward pass for a max pooling layer in the function `max_pool_backward_naive` in the file `cs231n/layers.py`. As always we check the correctness of the backward pass using numerical gradient checking.

```
[33]: x = np.random.randn(3, 2, 8, 8)
      dout = np.random.randn(3, 2, 4, 4)
      pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

      dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(x,
      ↪pool_param)[0], x, dout)

      out, cache = max_pool_forward_naive(x, pool_param)
      dx = max_pool_backward_naive(dout, cache)

      # Your error should be around 1e-12
      print('Testing max_pool_backward_naive function:')
      print('dx error: ', rel_error(dx, dx_num))
```

Testing max\_pool\_backward\_naive function:  
dx error: 3.275628496082982e-12

## 12 Fast layers

Making convolution and pooling layers fast can be challenging. To spare you the pain, we've provided fast implementations of the forward and backward passes for convolution and pooling layers in the file `cs231n/fast_layers.py`.

The fast convolution implementation depends on a Cython extension; to compile it you need to run the following from the `cs231n` directory:

```
python setup.py build_ext --inplace
```

The API for the fast versions of the convolution and pooling layers is exactly the same as the naive versions that you implemented above: the forward pass receives data, weights, and parameters and produces outputs and a cache object; the backward pass receives upstream derivatives and the cache object and produces gradients with respect to the data and weights.

**NOTE:** The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the following:

```
[30]: from cs231n.fast_layers import conv_forward_fast, conv_backward_fast
      from time import time

      x = np.random.randn(100, 3, 31, 31)
      w = np.random.randn(25, 3, 3, 3)
      b = np.random.randn(25,)
      dout = np.random.randn(100, 25, 16, 16)
      conv_param = {'stride': 2, 'pad': 1}

      t0 = time()
      out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
      t1 = time()
      out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
      t2 = time()

      print('Testing conv_forward_fast:')
      print('Naive: %fs' % (t1 - t0))
      print('Fast: %fs' % (t2 - t1))
      print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
      print('Difference: ', rel_error(out_naive, out_fast))

      t0 = time()
      dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
      t1 = time()
      dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
      t2 = time()

      print('\nTesting conv_backward_fast:')
      print('Naive: %fs' % (t1 - t0))
      print('Fast: %fs' % (t2 - t1))
      print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
      print('dx difference: ', rel_error(dx_naive, dx_fast))
      print('dw difference: ', rel_error(dw_naive, dw_fast))
      print('db difference: ', rel_error(db_naive, db_fast))
```

```
Testing conv_forward_fast:
Naive: 3.885641s
Fast: 0.007948s
Speedup: 488.872994x
Difference: 3.978240059466707e-11
```

```
Testing conv_backward_fast:
Naive: 14.729645s
Fast: 0.008946s
Speedup: 1646.516977x
dx difference: 1.2479237101602112e-11
dw difference: 1.5251920225304764e-12
```

db difference: 0.0

```
[34]: from cs231n.fast_layers import max_pool_forward_fast, max_pool_backward_fast

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
```

Testing pool\_forward\_fast:

Naive: 0.141594s

fast: 0.002992s

speedup: 47.321753x

difference: 0.0

Testing pool\_backward\_fast:

Naive: 0.387990s

speedup: 43.224229x

dx difference: 0.0

## 13 Sandwich layers

There are a couple common layer “sandwiches” that frequently appear in ConvNets. For example convolutional layers are frequently followed by ReLU and pooling, and affine layers are frequently

followed by ReLU. To make it more convenient to use these common patterns, we have defined several convenience layers in the file `cs231n/layer_utils.py`. Lets grad-check them to make sure that they work correctly:

```
[11]: from cs231n.layer_utils import conv_relu_pool_forward, conv_relu_pool_backward

x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(x, w,
    ↪b, conv_param, pool_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(x, w,
    ↪b, conv_param, pool_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(x, w,
    ↪b, conv_param, pool_param)[0], b, dout)

print('Testing conv_relu_pool_forward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing conv_relu_pool_forward:
dx error:  2.2041086893381568e-07
dw error:  2.717511280801651e-10
db error:  3.502737338077941e-11
```

```
[12]: from cs231n.layer_utils import conv_relu_forward, conv_relu_backward

x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w, b,
    ↪conv_param)[0], x, dout)
```

```

dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w, b,
    ↪conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w, b,
    ↪conv_param)[0], b, dout)

print('Testing conv_relu_forward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

```

Testing conv\_relu\_forward:  
dx error: 2.6296328828998218e-08  
dw error: 1.5342896567838248e-09  
db error: 7.688682989110673e-12

```

[13]: from cs231n.layer_utils import affine_relu_forward, affine_relu_backward

x = np.random.randn(2, 3, 4)
w = np.random.randn(12, 10)
b = np.random.randn(10)
dout = np.random.randn(2, 10)

out, cache = affine_relu_forward(x, w, b)
dx, dw, db = affine_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: affine_relu_forward(x, w,
    ↪b)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: affine_relu_forward(x, w,
    ↪b)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: affine_relu_forward(x, w,
    ↪b)[0], b, dout)

print('Testing affine_relu_forward:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

```

Testing affine\_relu\_forward:  
dx error: 4.5035126837662996e-09  
dw error: 6.687598198525906e-11  
db error: 3.2755814073540085e-12

[ ]:

# convnet

February 10, 2020

## 1 Train a ConvNet!

We now have a generic solver and a bunch of modularized layers. It's time to put it all together, and train a ConvNet to recognize the classes in CIFAR-10. In this notebook we will walk you through training a simple two-layer ConvNet and then set you free to build the best net that you can to perform well on CIFAR-10.

Open up the file `cs231n/classifiers/convnet.py`; you will see that the `two_layer_convnet` function computes the loss and gradients for a two-layer ConvNet. Note that this function uses the “sandwich” layers defined in `cs231n/layer_utils.py`.

```
[1]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from cs231n.classifier_trainer import ClassifierTrainer
from cs231n.gradient_check import eval_numerical_gradient
from cs231n.classifiers.convnet import *

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/
# ↪ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

[9]: from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
```



*Load the CIFAR-10 dataset from disk and perform preprocessing to prepare it for the two-layer neural net classifier. These are the same steps as we used for the SVM, but condensed to a single function.*

"""

*# Load the raw CIFAR-10 data*

cifar10\_dir = 'cs231n/datasets/cifar-10-batches-py'

X\_train, y\_train, X\_test, y\_test = load\_CIFAR10(cifar10\_dir)

*# Subsample the data*

mask = range(num\_training, num\_training + num\_validation)

X\_val = X\_train[mask]

y\_val = y\_train[mask]

mask = range(num\_training)

X\_train = X\_train[mask]

y\_train = y\_train[mask]

mask = range(num\_test)

X\_test = X\_test[mask]

y\_test = y\_test[mask]

*# Normalize the data: subtract the mean image*

mean\_image = np.mean(X\_train, axis=0)

X\_train -= mean\_image

X\_val -= mean\_image

X\_test -= mean\_image

*# Transpose so that channels come first*

X\_train = X\_train.transpose(0, 3, 1, 2).copy()

X\_val = X\_val.transpose(0, 3, 1, 2).copy()

x\_test = X\_test.transpose(0, 3, 1, 2).copy()

return X\_train, y\_train, X\_val, y\_val, X\_test, y\_test

*# Invoke the above function to get our data.*

X\_train, y\_train, X\_val, y\_val, X\_test, y\_test = get\_CIFAR10\_data()

print('Train data shape: ', X\_train.shape)

print('Train labels shape: ', y\_train.shape)

print('Validation data shape: ', X\_val.shape)

print('Validation labels shape: ', y\_val.shape)

print('Test data shape: ', X\_test.shape)

print('Test labels shape: ', y\_test.shape)

Train data shape: (49000, 3, 32, 32)

Train labels shape: (49000,)

Validation data shape: (1000, 3, 32, 32)

Validation labels shape: (1000,)

Test data shape: (1000, 32, 32, 3)

Test labels shape: (1000,)

## 2 Sanity check loss

After you build a new network, one of the first things you should do is sanity check the loss. When we use the softmax loss, we expect the loss for random weights (and no regularization) to be about  $\log(C)$  for  $C$  classes. When we add regularization this should go up.

```
[10]: model = init_two_layer_convnet()

X = np.random.randn(100, 3, 32, 32)
y = np.random.randint(10, size=100)

loss, _ = two_layer_convnet(X, model, y, reg=0)

# Sanity check: Loss should be about log(10) = 2.3026
print('Sanity check loss (no regularization): ', loss)

# Sanity check: Loss should go up when you add regularization
loss, _ = two_layer_convnet(X, model, y, reg=1)
print('Sanity check loss (with regularization): ', loss)
```

Sanity check loss (no regularization): 2.3026302984328333

Sanity check loss (with regularization): 2.344375763367764

## 3 Gradient check

After the loss looks reasonable, you should always use numeric gradient checking to make sure that your backward pass is correct. When you use numeric gradient checking you should use a small amount of artificial data and a small number of neurons at each layer.

```
[11]: num_inputs = 2
input_shape = (3, 16, 16)
reg = 0.0
num_classes = 10
X = np.random.randn(num_inputs, *input_shape)
y = np.random.randint(num_classes, size=num_inputs)

model = init_two_layer_convnet(num_filters=3, filter_size=3,
    ↪input_shape=input_shape)
loss, grads = two_layer_convnet(X, model, y)
for param_name in sorted(grads):
    f = lambda _: two_layer_convnet(X, model, y)[0]
    param_grad_num = eval_numerical_gradient(f, model[param_name],
    ↪verbose=False, h=1e-6)
```

```

    e = rel_error(param_grad_num, grads[param_name])
    print('%s max relative error: %e' % (param_name, rel_error(param_grad_num,
↳ grads[param_name])))

```

```

W1 max relative error: 3.023856e-07
W2 max relative error: 1.418324e-05
b1 max relative error: 2.668192e-08
b2 max relative error: 1.995789e-09

```

## 4 Overfit small data

A nice trick is to train your model with just a few training samples. You should be able to overfit small datasets, which will result in very high training accuracy and comparatively low validation accuracy.

```

[13]: # Use a two-layer ConvNet to overfit 50 training examples.

model = init_two_layer_convnet()
trainer = ClassifierTrainer()
best_model, loss_history, train_acc_history, val_acc_history = trainer.train(
    X_train[:50], y_train[:50], X_val, y_val, model, two_layer_convnet,
    reg=0.05, momentum=0.9, learning_rate=0.00001, batch_size=10,
↳ num_epochs=10,
    verbose=True)

```

```

starting iteration  0
Finished epoch 0 / 10: cost 2.304627, train: 0.100000, val 0.110000, lr
1.000000e-05
Finished epoch 1 / 10: cost 2.305375, train: 0.200000, val 0.110000, lr
9.500000e-06
Finished epoch 2 / 10: cost 2.273784, train: 0.320000, val 0.137000, lr
9.025000e-06
starting iteration  10
Finished epoch 3 / 10: cost 2.279856, train: 0.300000, val 0.138000, lr
8.573750e-06
Finished epoch 4 / 10: cost 2.231050, train: 0.340000, val 0.147000, lr
8.145063e-06
starting iteration  20
Finished epoch 5 / 10: cost 2.107130, train: 0.320000, val 0.150000, lr
7.737809e-06
Finished epoch 6 / 10: cost 2.040741, train: 0.300000, val 0.154000, lr
7.350919e-06
starting iteration  30
Finished epoch 7 / 10: cost 2.194671, train: 0.300000, val 0.152000, lr
6.983373e-06
Finished epoch 8 / 10: cost 2.082568, train: 0.320000, val 0.148000, lr

```

6.634204e-06

starting iteration 40

Finished epoch 9 / 10: cost 2.090183, train: 0.300000, val 0.148000, lr

6.302494e-06

Finished epoch 10 / 10: cost 1.967252, train: 0.320000, val 0.142000, lr


5.987369e-06

finished optimization. best validation accuracy: 0.154000

Plotting the loss, training accuracy, and validation accuracy should show clear overfitting:

```
[14]: plt.subplot(2, 1, 1)
plt.plot(loss_history)
plt.xlabel('iteration')
plt.ylabel('loss')

plt.subplot(2, 1, 2)
plt.plot(train_acc_history)
plt.plot(val_acc_history)
plt.legend(['train', 'val'], loc='upper left')
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.show()
```



## 5 Train the net

Once the above works, training the net is the next thing to try. You can set the `acc_frequency` parameter to change the frequency at which the training and validation set accuracies are tested. If your parameters are set properly, you should see the training and validation accuracy start to improve within a hundred iterations, and you should be able to train a reasonable model with just one epoch.

Using the parameters below you should be able to get around 50% accuracy on the validation set.

```
[15]: model = init_two_layer_convnet(filter_size=7)
      trainer = ClassifierTrainer()
      best_model, loss_history, train_acc_history, val_acc_history = trainer.train(
          X_train, y_train, X_val, y_val, model, two_layer_convnet,
          reg=0.001, momentum=0.9, learning_rate=0.0001, batch_size=50,
      ↪ num_epochs=1,
          acc_frequency=50, verbose=True)
```

```
starting iteration  0
Finished epoch 0 / 1: cost 2.294376, train: 0.105000, val 0.108000, lr
1.000000e-04
starting iteration  10
starting iteration  20
starting iteration  30
starting iteration  40
starting iteration  50
Finished epoch 0 / 1: cost 1.934034, train: 0.301000, val 0.320000, lr
1.000000e-04
starting iteration  60
starting iteration  70
starting iteration  80
starting iteration  90
starting iteration 100
Finished epoch 0 / 1: cost 1.768556, train: 0.373000, val 0.355000, lr
1.000000e-04
starting iteration 110
starting iteration 120
starting iteration 130
starting iteration 140
starting iteration 150
Finished epoch 0 / 1: cost 1.515819, train: 0.381000, val 0.378000, lr
1.000000e-04
starting iteration 160
starting iteration 170
```

starting iteration 180  
starting iteration 190  
starting iteration 200  
Finished epoch 0 / 1: cost 1.677729, train: 0.375000, val 0.411000, lr  
1.000000e-04  
starting iteration 210  
starting iteration 220  
starting iteration 230  
starting iteration 240  
starting iteration 250  
Finished epoch 0 / 1: cost 1.849573, train: 0.388000, val 0.413000, lr  
1.000000e-04  
starting iteration 260  
starting iteration 270  
starting iteration 280  
starting iteration 290  
starting iteration 300  
Finished epoch 0 / 1: cost 1.854043, train: 0.429000, val 0.426000, lr  
1.000000e-04  
starting iteration 310  
starting iteration 320  
starting iteration 330  
starting iteration 340  
starting iteration 350  
Finished epoch 0 / 1: cost 1.815773, train: 0.456000, val 0.441000, lr  
1.000000e-04  
starting iteration 360  
starting iteration 370  
starting iteration 380  
starting iteration 390  
starting iteration 400  
Finished epoch 0 / 1: cost 1.995193, train: 0.424000, val 0.427000, lr  
1.000000e-04  
starting iteration 410  
starting iteration 420  
starting iteration 430  
starting iteration 440  
starting iteration 450  
Finished epoch 0 / 1: cost 2.085419, train: 0.426000, val 0.436000, lr  
1.000000e-04  
starting iteration 460  
starting iteration 470  
starting iteration 480  
starting iteration 490  
starting iteration 500  
Finished epoch 0 / 1: cost 1.864259, train: 0.441000, val 0.432000, lr  
1.000000e-04  
starting iteration 510

starting iteration 520  
starting iteration 530  
starting iteration 540  
starting iteration 550  
Finished epoch 0 / 1: cost 1.329013, train: 0.484000, val 0.436000, lr 1.000000e-04  
starting iteration 560  
starting iteration 570  
starting iteration 580  
starting iteration 590  
starting iteration 600  
Finished epoch 0 / 1: cost 1.509924, train: 0.464000, val 0.473000, lr 1.000000e-04  
starting iteration 610  
starting iteration 620  
starting iteration 630  
starting iteration 640  
starting iteration 650  
Finished epoch 0 / 1: cost 1.282254, train: 0.454000, val 0.472000, lr 1.000000e-04  
starting iteration 660  
starting iteration 670  
starting iteration 680  
starting iteration 690  
starting iteration 700  
Finished epoch 0 / 1: cost 1.523986, train: 0.464000, val 0.436000, lr 1.000000e-04  
starting iteration 710  
starting iteration 720  
starting iteration 730  
starting iteration 740  
starting iteration 750  
Finished epoch 0 / 1: cost 2.197756, train: 0.484000, val 0.469000, lr 1.000000e-04  
starting iteration 760  
starting iteration 770  
starting iteration 780  
starting iteration 790  
starting iteration 800  
Finished epoch 0 / 1: cost 1.176381, train: 0.495000, val 0.479000, lr 1.000000e-04  
starting iteration 810  
starting iteration 820  
starting iteration 830  
starting iteration 840  
starting iteration 850  
Finished epoch 0 / 1: cost 1.670690, train: 0.478000, val 0.488000, lr 1.000000e-04

```
starting iteration 860
starting iteration 870
starting iteration 880
starting iteration 890
starting iteration 900
Finished epoch 0 / 1: cost 1.255807, train: 0.451000, val 0.476000, lr
1.000000e-04
starting iteration 910
starting iteration 920
starting iteration 930
starting iteration 940
starting iteration 950
Finished epoch 0 / 1: cost 1.237256, train: 0.470000, val 0.462000, lr
1.000000e-04
starting iteration 960
starting iteration 970
Finished epoch 1 / 1: cost 1.929538, train: 0.485000, val 0.481000, lr
9.500000e-05
finished optimization. best validation accuracy: 0.488000
```

## 6 Visualize weights


We can visualize the convolutional weights from the first layer. If everything worked properly, these will usually be edges and blobs of various colors and orientations.

```
[16]: from cs231n.vis_utils import visualize_grid

      grid = visualize_grid(best_model['W1'].transpose(0, 2, 3, 1))
      plt.imshow(grid.astype('uint8'))
```




```
[16]: <matplotlib.image.AxesImage at 0x1bb4e1c3ec8>
```






[ ]:


Softmax






Softmax Learning Curve




Softmax Validation Accuracy During Training




Twolayernn:






Twolayernn Learning Curve




Twolayernn Validation Accuracy During Training




Convnet:



Convnet Learning Curve



Convnet Validation Accuracy During Training



Mymodel:

